

Chapter 2 - Installation and Use of NetX Duo

Contents

Chapter 2 - Installation and Use of NetX Duo	1
Host Considerations	1
Target Considerations	2
Product Distribution	2
NetX Duo Installation	2
Using NetX Duo	3
Troubleshooting	3
Configuration Options	4
System Configuration Options	4
ARP Configuration Options	6
ICMP Configuration Options	8
IGMP Configuration Options	9
IP Configuration Options	9
Packet Configuration Options	11
RARP Configuration Options	13
TCP Configuration Options	13
UDP Configuration Options	18
IPv6 Options	19
Neighbor Cache Configuration Options	21
Miscellaneous ICMPv6 Configuration Options	23
NetX Duo Version ID	25

Chapter 2 - Installation and Use of NetX Duo

This chapter contains a description of various issues related to installation, setup, and use of the high- performance network stack NetX Duo, including the following:

Host Considerations

Embedded development is usually performed on Windows or Linux (Unix) host computers. After the application is compiled, linked, and the executable is generated on the host, it is downloaded to the target hardware for execution.

Usually the target download is done from within the development tool's debugger. After download, the debugger is responsible for providing target execution control (go, halt, breakpoint, etc.) as well as access to memory and processor registers.

Most development tool debuggers communicate with the target hardware via on-chip debug (OCD) connections such as JTAG (IEEE 1149.1) and Background Debug Mode (BDM). Debuggers also communicate with target hardware through In-Circuit Emulation (ICE) connections. Both OCD and ICE connections provide robust solutions with minimal intrusion on the target resident software.

As for resources used on the host, the source code for NetX Duo is delivered in ASCII format and requires approximately 1 Mbytes of space on the host computer's hard disk.

Target Considerations

NetX Duo requires between 5 KBytes and 45 KBytes of Read-Only Memory (ROM) on the target. Another 1 to 5KBytes of the target's Random Access Memory (RAM) are required for the NetX Duo thread stack and other global data structures.

In addition, NetX Duo requires the use of two ThreadX timer objects and one ThreadX mutex object. These facilities are used for periodic processing needs and thread protection inside the NetX Duo protocol stack.

Product Distribution

NetX Duo can be obtained from our public source code repository at <https://github.com/eclipse-threadx/netxduo/>.

The following is a list of several important files in the repository:

nx_api.h

C header file containing all system equates, data structures, and service prototypes.

nx_port.h C header file containing all development-tool and targetspecific data definitions and structures.

demo_netx.c C file containing a small demo application.

nx.a (or nx.lib)

Binary version of the NetX Duo C library that is distributed with the standard package.

NetX Duo Installation

NetX Duo is installed by cloning the GitHub repository to your local machine. The following is typical syntax for creating a clone of the NetX Duo repository on your PC:

```
git clone https://github.com/eclipse-threadx/netxduo
```

Alternatively you can download a copy of the repository using the download button on the GitHub main page.

You will also find instructions for building the NetX Duo library on the front page of the online repository.

Important: *Application software needs access to the NetX Duo library file (usually ***nx.a*** or ***nx.lib***) and the C include files ***nx_api.h***, and ***nx_port.h***. This is accomplished either by setting the appropriate path for the development tools or by copying these files into the application development area.*

Using NetX Duo

Using NetX Duo is easy. Basically, the application code must include ***nx_api.h*** during compilation and link with the NetX Duo library ***nx.a*** (or ***nx.lib***).

The following are the four easy steps required to build a NetX Duo application:

Step	Description
Step 1:	Include the <i>nx_api.h</i> file in all application files that use NetX Duo services or data structures.
Step 2:	Initialize the NetX Duo system by calling <i>nx_system_initialize</i> from the <i>tx_application_define</i> function or an application thread.
Step 3:	Create an IP instance, enable the Address Resolution Protocol (ARP), if necessary, and any sockets after <i>nx_system_initialize</i> is called.
Step 4:	Compile application source and link with the NetX Duo runtime library <i>nx.a</i> (or <i>nx.lib</i>). The resulting image can be downloaded to the target and executed!

Troubleshooting

Each NetX Duo port is delivered with one or more demonstrations that execute on an actual network or via a simulated network driver. It is always a good idea to get the demonstration system running first.

If the demonstration system does not run properly, perform the following operations to narrow the problem:

1. Determine how much of the demonstration is running.
2. Increase stack sizes in any new application threads.
3. Recompile the NetX Duo library with the appropriate debug options listed in the configuration option section.
4. Examine the NX_IP structure to see if packets are being sent or received.
5. Examine the default packet pool to see if there are available packets.
6. Ensure the network driver is supplying ARP and IP packets with its headers on 4-byte boundaries for applications requiring IPv4 or IPv6 connectivity.
7. Temporarily bypass any recent changes to see if the problem disappears or changes.

Configuration Options

There are several configuration options when building the NetX Duo library and the application using NetX Duo. The configuration options can be defined in the application source, on the command line, or within the *nx_user.h* include file, unless otherwise specified.

Note: *Options defined in nx_user.h are applied only if the application and NetX Duo library are built with NX_INCLUDE_USER_DEFINE_FILE defined.*

The following sections list the configuration options available in NetX Duo. General options applicable to both IPv4 and IPv6 are listed first, followed by IPv6-specific options.

System Configuration Options

Option	Description
NX_ASSERT_FAIL	Symbol that defines the debug statement to use when an assertion fails.
NX_DEBUG	Defined, enables the optional print debug information available from the RAM Ethernet network driver.
NX_DEBUG_PACKET	Defined, enables the optional debug packet dumping available in the RAM Ethernet network driver.
NX_DISABLE_ASSERT	Defined, disables ASSERT checks in the source code. By default this option is not defined.

Option	Description
NX_DISABLE_ERROR_CHECKING	Defined, removes the basic NetX Duo error checking API and improves performance. API return codes not affected by disabling error checking are listed in bold typeface in the API definition. This define is typically used after the application is debugged sufficiently and its use improves performance and decreases code size.
NX_DRIVER_DEFERRED_PROCESSING	Defined, enables deferred network driver packet handling. This allows the network driver to place a packet on the IP instance and have the real processing routine called from the NetX Duo internal IP helper thread.
NX_DUAL_PACKET_POOL_ENABLE	Renamed to NX_ENABLE_DUAL_PACKET_POOL . Although it is still being supported, new designs are encouraged to use NX_ENABLE_DUAL_PACKET_POOL .
NX_ENABLE_DUAL_PACKET_POOL	Defined, allows the stack to use two packet pools, one with large payload size and one with smaller payload size. By default this option is not enabled.
NX_ENABLE_EXTENDED_NOTIFY_SUPPORT	Supports more callback hooks in the stack. These callback functions are used by the BSD wrapper layer. By default this option is not defined.
NX_ENABLE_INTERFACE_CAPABILITY	Defined, allows the interface device driver to specify extra capability information, such as checksum off-loading. By default this option is not defined.
NX_ENABLE_SOURCE_ADDRESS_CHECK	Defined, enables the source address of incoming packet to be checked. By default this option is disabled.
NX_IPSEC_ENABLE	Defined, enables the NetX Duo library to support IPsec operations. This feature requires the optional NetX Duo IPsec module. By default this feature is not enabled.

Option	Description
NX_LITTLE_ENDIAN	Defined, performs the necessary byte swapping on little endian environments to ensure the protocol headers are in proper big endian format. Note the default is typically setup in <i>nx_port.h</i> .
NX_MAX_PHYSICAL_INTERFACES	Specifies the total number of physical network interfaces on the device. The default value is 1 and is defined in <i>nx_api.h</i> ; a device must have at least one physical interface. Note this does not include the loopback interface.
NX_NAT_ENABLE	Defined, NetX Duo is built with NAT process. By default this option is not defined.
NX_PHYSICAL_HEADER	Specifies the size in bytes of the physical header of the frame. The default value is 16 (based on a typical 14-byte Ethernet frame aligned to 32-bit boundary) and is defined in <i>nx_api.h</i> . The application can override the default by defining the value before <i>nx_api.h</i> is included, such as in <i>nx_user.h</i> .
NX_PHYSICAL_TRAILER	Specifies the size in bytes of the physical packet trailer and is typically used to reserve storage for things like Ethernet CRCs, etc. The default value is 4 and is defined in <i>nx_api.h</i> .

ARP Configuration Options

Option	Description
NX_ARP_DEFEND_BY_REPLY	Defined, allows NetX Duo to defend its IP address by sending an ARP response.
NX_ARP_DEFEND_INTERVAL	Defines the interval, in seconds, the ARP module sends out the next defend packet in response to an incoming ARP message that indicates an address in conflict.

Option	Description
NX_ARP_DISABLE_AUTO_ARP_ENTRY	Defined to NX_DISABLE_ARP_AUTO_ENTRY . Although it is still being supported, new designs are encouraged to use NX_DISABLE_ARP_AUTO_ENTRY .
NX_ARP_EXPIRATION_RATE	Specifies the number of seconds ARP entries remain valid. The default value of zero disables expiration or aging of ARP entries and is defined in <i>nx_api.h</i> . The application can override the default by defining the value before <i>nx_api.h</i> is included.
NX_ARP_MAC_CHANGE_NOTIFICATION_ENABLE	Defined to NX_ENABLE_ARP_MAC_CHANGE_NOTIFICATION . Although it is still being supported, new designs are encouraged to use NX_ENABLE_ARP_MAC_CHANGE_NOTIFICATION .
NX_ARP_MAX_QUEUE_DEPTH	Specifies the maximum number of packets that can be queued while waiting for an ARP response. The default value is 4 and is defined in <i>nx_api.h</i> .
NX_ARP_MAXIMUM_RETRIES	Specifies the maximum number of ARP retries made without an ARP response. The default value is 18 and is defined in <i>nx_api.h</i> . The application can override the default by defining the value before <i>nx_api.h</i> is included.
NX_ARP_UPDATE_RATE	Specifies the number of seconds between ARP retries. The default value is 10, which represents 10 seconds, and is defined in <i>nx_api.h</i> . The application can override the default by defining the value before <i>nx_api.h</i> is included.
NX_DISABLE_ARP_AUTO_ENTRY	Defined, disables entering ARP request information in the ARP cache.
NX_DISABLE_ARP_INFO	Defined, disables ARP information gathering.
NX_ENABLE_ARP_MAC_CHANGE_NOTIFICATION	Defined to NX_ENABLE_ARP_MAC_CHANGE_NOTIFICATION ARP to invoke a callback notify function on detecting the MAC address is updated.

ICMP Configuration Options

Option	Description
NX_DISABLE_ICMP_INFO	Defined, disables ICMP information gathering.
NX_DISABLE_ICMP_RX_CHECKSUM	Defined, disables both ICMPv4 and ICMPv6 checksum computation on received ICMP packets. This option is useful when the network interface driver is able to verify the ICMPv4 and ICMPv6 checksum, and the application does not use the IP fragmentation feature or the IPsec feature. By default this option is not defined.
NX_DISABLE_ICMP_TX_CHECKSUM	Defined, disables both ICMPv4 and ICMPv6 checksum computation on transmitted ICMP packets. This option is useful where the network interface driver is able to compute the ICMPv4 and ICMPv6 checksum, and the application does not use the IP fragmentation feature or IPsec feature. By default this option is not defined.
NX_DISABLE_ICMPV4_ERROR_MESSAGE	Defined, NetX Duo does not send ICMPv4 Error Messages in response to error conditions such as improperly formatted IPv4 header. By default this option is not defined.
NX_DISABLE_ICMPV4_RX_CHECKSUM	Defined, disables ICMPv4 checksum computation on received ICMP packets. This option is defined automatically if NX_DISABLE_ICMP_RX_CHECKSUM is defined. By default this option is not defined.
NX_DISABLE_ICMPV4_RX_CHECKSUM	Renamed to NX_DISABLE_ICMPV4_RX_CHECKSUM . Although it is still being supported, new designs are encouraged to use NX_DISABLE_ICMPV4_RX_CHECKSUM .

Option	Description
NX_DISABLE_ICMPV4_TX_CHECKSUM	Defined, disables ICMPv4 checksum computation on transmitted ICMP packets. This option is defined automatically if NX_DISABLE_ICMP_TX_CHECKSUM is defined. By default this option is not defined.
NX_DISABLE_ICMPv4_TX_CHECKSUM	Moved to NX_DISABLE_ICMPV4_TX_CHECKSUM . Although it is still being supported, new designs are encouraged to use NX_DISABLE_ICMPV4_TX_CHECKSUM .
NX_ENABLE_ICMP_ADDRESS_CHECK	Defined, the destination address of ICMP packet is checked. The default is disabled. An ICMP Echo Request destined to an IP broadcast or IP multicast address will be silently discarded.

IGMP Configuration Options

Option	Description
NX_DISABLE_IGMP_INFO	Defined, disables IGMP information gathering.
NX_DISABLE_IGMPV2	Defined, disables IGMPv2 support, and NetX Duo supports IGMPv1 only. By default this option is not set and is defined in nx_api.h .
NX_MAX_MULTICAST_GROUPS	Specifies the maximum number of multicast groups that can be joined. The default value is 7 and is defined in nx_api.h . The application can override the default by defining the value before nx_api.h is included.

IP Configuration Options

Option	Description
NX_DISABLE_FRAGMENTATION	Defined, disables both IPv4 and IPv6 fragmentation and reassembly logic.

Option	Description
NX_DISABLE_IPV4	Defined, disables IPv4 functionality. This option can be used to build NetX Duo to support IPv6 only. By default this option is not defined.
NX_DISABLE_IP_INFO	Defined, disables IP information gathering.
NX_DISABLE_IP_RX_CHECKSUM	Defined, disables checksum logic on received IPv4 packets. This is useful if the network device is able to verify the IPv4 checksum, and the application does not expect to use IP fragmentation or IPsec.
NX_DISABLE_IP_TX_CHECKSUM	Defined, disables checksum logic on IPv4 packets sent. This is useful in situations in which the underlying network device is capable of generating the IPv4 header checksum, and the application does not expect to use IP fragmentation or IPsec.
NX_DISABLE_LOOPBACK_INTERFACE	Defined, disables NetX Duo support for the loopback interface.
NX_DISABLE_RX_SIZE_CHECKING	Defined, disables the size checking on received packets.
NX_ENABLE_IP_RAW_PACKET_FILTER	Defined, enables the IP raw packet receive filter functionality. Applications requiring more control over the type of raw IP packets to be received can use this feature. The IP raw packet filter feature also supports the raw socket operation in the BSD compatibility layer. By default this option is not defined.
NX_ENABLE_IP_STATIC_ROUTING	Defined, enables IPv4 static routing in which a destination address can be assigned a specific next hop address. By default IPv4 static routing is disabled.
NX_FRAGMENT_IMMEDIATE_ASSEMBLY	Defined, allows IPv4 and IPv6 reassembly logic to execute right away after receiving an IP fragment. By default this option is not defined.

Option	Description
NX_IP_MAX_REASSEMBLY_TIME	Symbol that controls maximum time allowed to reassemble IPv4 fragment and IPv6 fragment. Note the value defined here overwrites both NX_IPV4_MAX_REASSEMBLY_TIME and NX_IPV6_MAX_REASSEMBLY_TIME .
NX_IP_PERIODIC_RATE	Defined, specifies the number of ThreadX timer ticks in one second. The default value is derived from the ThreadX symbol TX_TIMER_TICKS_PER_SECOND , which by default is set to 100 (10ms timer). Applications shall exercise caution when modifying this value, as the rest of the NetX Duo modules derive timing information from NX_IP_PERIODIC_RATE .
NX_IP_RAW_MAX_QUEUE_DEPTH	Symbol that controls the number of raw IP packets can be queued on the raw packet receive queue. By default value is set to 20.
NX_IP_ROUTING_TABLE_SIZE	Defined, sets the maximum number of entries in the IPv4 static routing table, which is a list of an outgoing interface and the next hop addresses for a given destination address. The default value is 8 and is defined in nx_api.h . This symbol is used only if NX_ENABLE_IP_STATIC_ROUTING is defined.
NX_IPV4_MAX_REASSEMBLY_TIME	Symbol that controls maximum time allowed to reassemble IPv4 fragment. Note the value defined in NX_IP_MAX_REASSEMBLY_TIME overwrites this value.
NX_ENABLE_TCPIP_OFFLOAD	Symbol that enables TCP/IP offload feature. Note NX_ENABLE_INTERFACE_CAPABILITY must be defined to enable this feature.

Packet Configuration Options

Option	Description
NX_DISABLE_PACKET_CHAIN	Defined, disables the packet chain logic. By default this is not defined.
NX_DISABLE_PACKET_INFO	Defined, disables packet pool information gathering.
NX_ENABLE_LOW_WATERMARK	Defined, enables NetX Duo packet pool low watermark feature. Application sets low watermark value. On receiving TCP packets, if the packet pool low watermark is reached, NetX Duo silently discards the packet by releasing it, preventing the packet pool from starvation. By default this feature is not enabled.
NX_ENABLE_PACKET_DEBUG_INFO	Defined, logs packet debug information.
NX_PACKET_ALIGNMENT	Defined, specifies the alignment requirement, in bytes, for starting address of the packet payload area. This option deprecates NX_PACKET_HEADER_PAD and NX_PACKET_HEADER_PAD_SIZE . By default this option is defined to be 4, making the starting address of the payload area 4-byte aligned.
NX_PACKET_HEADER_PAD	Defined, enables padding towards the end of the NX_PACKET control block. The number of ULONG words to pad is defined by NX_PACKET_HEADER_PAD_SIZE . Note this option is deprecated by NX_PACKET_ALIGNMENT .

Option	Description
NX_PACKET_HEADER_PAD_SIZE	Sets the number of ULONG words to be padded to the NX_PACKET structure, allowing the packet payload area to start at the desired alignment. This feature is useful when receive buffer descriptors point directly into NX_PACKET payload area, and the network interface receive logic or the cache operation logic expects the buffer starting address to meet certain alignment requirements. This value becomes valid only when NX_PACKET_HEADER_PAD is defined. Note this option is deprecated by NX_PACKET_ALIGNMENT .

RARP Configuration Options

Option	Description
NX_DISABLE_RARP_INFO	Defined, disables RARP information gathering.

TCP Configuration Options

Option	Description
NX_DISABLE_RESET_DISCONNECT	Defined, disables the reset processing during disconnect when the timeout value supplied is specified as NX_NO_WAIT .
NX_DISABLE_TCP_INFO	Defined, disables TCP information gathering.
NX_DISABLE_TCP_RX_CHECKSUM	Defined, disables checksum logic on received TCP packets. This is only useful in situations in which the link-layer has reliable checksum or CRC processing, or the interface driver is able to verify the TCP checksum in hardware, and the application does not use IPsec.

Option	Description
NX_DISABLE_TCP_TX_CHECKSUM	Defined, disables checksum logic for sending TCP packets. This is only useful in situations in which the receiving network node has received TCP checksum logic disabled or the underlying network driver is capable of generating the TCP checksum, and the application does not use IPsec.
NX_ENABLE_TCP_KEEPALIVE	Defined, enables the optional TCP keepalive timer. The default settings is not enabled.
NX_ENABLE_TCP_MSS_CHECK	Defined, enables the verification of minimum peer MSS before accepting a TCP connection. To use this feature, the symbol NX_ENABLE_TCP_MSS_MINIMUM must be defined. By default, this option is not enabled.
NX_ENABLE_TCP_QUEUE_DEPTH_NOTIFY	Defined, application to install a callback function that is invoked when the TCP transmit queue depth is no longer at maximum value. This callback serves as an indication that the TCP socket is ready to transmit more data. By default this option is not enabled.
NX_ENABLE_TCP_WINDOW_SCALE	Defines the window scaling option for TCP applications. If defined, window scaling option is negotiated during TCP connection phase, and the application is able to specify a window size larger than 64K. The default setting is not enabled (not defined).
NX_MAX_LISTEN_REQUESTS	Specifies the maximum number of server listen requests. The default value is 10 and is defined in nx_api.h . The application can override the default by defining the value before nx_api.h is included.

Option	Description
NX_TCP_ACK_EVERY_N_PACKETS	Specifies the number of TCP packets to receive before sending an ACK. Note if NX_TCP_IMMEDIATE_ACK is enabled but NX_TCP_ACK_EVERY_N_PACKETS is not, this value is automatically set to 1 for backward compatibility.
NX_TCP_ACK_TIMER_RATE	Specifies how the number of system ticks (NX_IP_PERIODIC_RATE) is divided to calculate the timer rate for the TCP delayed ACK processing. The default value is 5, which represents 200ms, and is defined in <i>nx_tcp.h</i> . The application can override the default by defining the value before <i>nx_api.h</i> is included. Renamed to NX_ENABLE_TCP_KEEPALIVE .
NX_TCP_ENABLE_KEEPALIVE	Although it is still being supported, new designs are encouraged to use NX_ENABLE_TCP_KEEPALIVE .
NX_TCP_ENABLE_MSS_CHECK	Renamed to NX_ENABLE_TCP_MSS_CHECK . Although it is still being supported, new designs are encouraged to use NX_ENABLE_TCP_MSS_CHECK .
NX_TCP_ENABLE_WINDOW_SCALING	Renamed to NX_ENABLE_TCP_WINDOW_SCALING . Although it is still being supported, new designs are encouraged to use NX_ENABLE_TCP_WINDOW_SCALING .

Option	Description
NX_TCP_FAST_TIMER_RATE	Specifies how the number of NetX Duo internal ticks (NX_IP_PERIODIC_RATE) is divided to calculate the fast TCP timer rate. The fast TCP timer is used to drive the various TCP timers, including the delayed ACK timer. The default value is 10, which represents 100ms assuming the ThreadX timer is running at 10ms. This value is defined in <i>nx_tcp.h</i> . The application can override the default by defining the value before <i>nx_api.h</i> is included.
NX_TCP_IMMEDIATE_ACK	Defined, enables the optional TCP immediate ACK response processing. Defining this symbol is equivalent to defining NX_TCP_ACK_EVERY_N_PACKETS to be 1.
NX_TCP_KEEPALIVE_INITIAL	Specifies the number of seconds of inactivity before the keepalive timer activates. The default value is 7200, which represents 2 hours, and is defined in <i>nx_tcp.h</i> . The application can override the default by defining the value before <i>nx_api.h</i> is included.
NX_TCP_KEEPALIVE_RETRIES	Specifies how many keepalive retries are allowed before the connection is deemed broken. The default value is 10, which represents 10 retries, and is defined in <i>nx_tcp.h</i> . The application can override the default by defining the value before <i>nx_api.h</i> is included.
NX_TCP_KEEPALIVE_RETRY	Specifies the number of seconds between retries of the keepalive timer assuming the other side of the connection is not responding. The default value is 75, which represents 75 seconds between retries, and is defined in <i>nx_tcp.h</i> . The application can override the default by defining the value before <i>nx_api.h</i> is included.

Option	Description
NX_TCP_MAX_OUT_OF_ORDER_SACKERS	Symbol that defines the maximum number of out-of-order TCP packets can be kept in the TCP socket receive queue. This symbol can be used to limit the number of packets queued in the TCP receive socket, preventing the packet pool from being starved. By default this symbol is not defined, thus there is no limit on the number of out of order packets being queued in the TCP socket.
NX_TCP_MAXIMUM_RETRIES	Specifies how many data transmit retries are allowed before the connection is deemed broken. The default value is 10, which represents 10 retries, and is defined in <i>nx_tcp.h</i> . The application can override the default by defining the value before <i>nx_api.h</i> is included.
NX_TCP_MAXIMUM_RX_QUEUE	Symbol that defines the maximum receive queue for TCP sockets. This feature is enabled by NX_ENABLE_LOW_WATERMARK .
NX_TCP_MAXIMUM_TX_QUEUE	Specifies the maximum depth of the TCP transmit queue before TCP send requests are suspended or rejected. The default value is 20, which means that a maximum of 20 packets can be in the transmit queue at any given time. Note packets stay in the transmit queue until an ACK that covers some or all of the packet data is received from the other side of the connection. This constant is defined in <i>nx_tcp.h</i> . The application can override the default by defining the value before <i>nx_api.h</i> is included.
NX_TCP_MSS_MINIMUM	Symbol that defines the minimal MSS value NetX Duo TCP module accepts. This feature is enabled by NX_ENABLE_TCP_MSS_CHECK .

Option	Description
NX_TCP_QUEUE_DEPTH_UPDATE_NOTIFY	<i>NX_ENABLE_TCP_QUEUE_DEPTH_UPDATE_NOTIFY</i> Although it is still being supported, new designs are encouraged to use <i>NX_ENABLE_TCP_QUEUE_DEPTH_UPDATE_NOTIFY</i> .
NX_TCP_RETRY_SHIFT	Specifies how the retransmit timeout period changes between retries. If this value is 0, the initial retransmit timeout is the same as subsequent retransmit timeouts. If this value is 1, each successive retransmit is twice as long. If this value is 2, each subsequent retransmit timeout is four times as long. The default value is 0 and is defined in <i>nx_tcp.h</i> . The application can override the default by defining the value before <i>nx_api.h</i> is included.
NX_TCP_TRANSMIT_TIMER_RATE	Specifies how the number of system ticks (<i>NX_IP_PERIODIC_RATE</i>) is divided to calculate the timer rate for the TCP transmit retry processing. The default value is 1, which represents 1 second, and is defined in <i>nx_tcp.h</i> . The application can override the default by defining the value before <i>nx_api.h</i> is included.

UDP Configuration Options

Option	Description
NX_DISABLE_UDP_INFO	Defined, disables UDP information gathering.
NX_DISABLE_UDP_RX_CHECKSUM	Defined, disables the UDP checksum computation on incoming UDP packets. This is useful if the network interface driver is able to verify UDP header checksum in hardware, and the application does not enable IPsec or IP fragmentation logic.

Option	Description
NX_DISABLE_UDP_TX_CHECKSUM	Defined, disables the UDP checksum computation on outgoing UDP packets. This is useful if the network interface driver is able to compute UDP header checksum and insert the value in the IP head before transmitting the data, and the application does not enable IPsec or IP fragmentation logic.

IPv6 Options

Option	Description
NX_DISABLE_IPV6	Disables IPv6 functionality when the NetX Duo library is built. For applications that do not need IPv6, this avoids pulling in code and additional storage space needed to support IPv6.
NX_DISABLE_IPV6_PATH_MTU_DISCOVERY	Defines path MTU discovery, which is used to determine the maximum MTU in the path to a target in the NetX Duo host destination table. This enables the NetX Duo host to send the largest possible packet that will not require fragmentation. By default, this option is defined (path MTU is disabled).
NX_ENABLE_IPV6_ADDRESS_CHANGE_CALLBACK	Defines a callback function to be invoked when the IPv6 address is changed. By default this option is not enabled.
NX_ENABLE_IPV6_MULTICAST	Defined, enables IPv6 multicast join/leave function. By default this option is not enabled.
NX_ENABLE_IPV6_PATH_MTU_DISCOVERY	Enables the IPv6 path MTU discovery feature. By default this option is not enabled.

Option	Description
NX_IPV6_ADDRESS_CHANGE_NOTIFY_ENABLE	NX_ENABLE_IPV6_ADDRESS_CHANGE_NOTIFY. Although it is still being supported, new designs are encouraged to use NX_ENABLE_IPV6_ADDRESS_CHANGE_NOTIFY.
NX_IPV6_DEFAULT_ROUTER_TABLE_SIZE	Specifies the number of entries in the IPv6 routing table. At least one entry is needed for the default router. Defined in <i>nx_api.h</i> , the default value is 8.
NX_IPV6_DESTINATION_TABLE_SIZE	Specifies the number of entries in the IPv6 destination table. This stores information about next hop addresses for IPv6 addresses. Defined in <i>nx_api.h</i> , the default value is 8.
NX_IPV6_MAX_REASSEMBLY_TIME	Symbol that controls the maximum time allowed to reassemble IPv6 fragment.
NX_IPV6_MULTICAST_ENABLE	Renamed to NX_ENABLE_IPV6_MULTICAST. Although it is still being supported, new designs are encouraged to use NX_ENABLE_IPV6_MULTICAST.
NX_IPV6_PREFIX_LIST_TABLE_SIZE	Specifies the size of the prefix table. Prefix information is obtained from router advertisements and is part of the IPv6 address configuration. Defined in <i>nx_api.h</i> , the default value is 8.
NX_IPV6_STATELESS_AUTOCONFIG_DISABLE	Defines NetX Duo to disable stateless address autoconfiguration feature. By default this option is not enabled.
NX_MAX_IPV6_ADDRESSES	Specifies the number of entries in the IPv6 address pool. During interface configuration, NetX Duo uses IPv6 entries from the pool. It is defaulted to (NX_MAX_PHYSICAL_INTERFACES * 3) to allow each interface to have at least one link local address and two global addresses. Note that all interfaces share the IPv6 address pool.

Option	Description
NX_PATH_MTU_INCREASE_WAIT_ \$INTERVAL	wait interval in timer ticks to reset the path MTU for a specific target in the destination table. If NX_DISABLE_IPV6_PATH_MTU_DISCOVERY is defined, defining this symbol has no effect.
NX_PATH_MTU_INCREASE_WAIT_ \$INTERVAL	specifies the wait interval (in seconds) to reset the path MTU value for a destination table entry. It is valid only if NX_ENABLE_IPV6_PATH_MTU_DISCOVERY is defined. By default this value is set to 600 (seconds).

Neighbor Cache Configuration Options

Option	Description
NX_DELAY_FIRST_PROBE_TIME	Specifies the delay in seconds before the first solicitation is sent out for a cache entry in the STALE state. Defined in nx_nd_cache.h , the default value is 5.
NX_DISABLE_IPV6_DAD	Defined, this option disables Duplicate Address Detection (DAD) during IPv6 address assignment. Addresses are set either by manual configuration or through Stateless Address Auto Configuration.
NX_DISABLE_IPV6_PURGE_UNUSED	Defaces Entries events NetX Duo from removing older cache table entries before their timeout expires to make room for new entries when the table is full. Static and router entries are never purged.

Option	Description
NX_IPV6_DAD_TRANSMITS	Specifies the number of Neighbor Solicitation messages to be sent before NetX Duo marks an interface address as valid. If NX_DISABLE_IPV6_DAD is defined (DAD disabled), setting this option has no effect. Alternatively, a value of zero (0) turns off DAD but leaves the DAD functionality in NetX Duo. Defined in <i>nx_api.h</i> , the default value is 3.
NX_IPV6_DISABLE_PURGE_UNUSEDCACHE_ENTRIES	NX_DISABLE_IPV6_PURGE_UNUSED_CACHE_ENTRIES Although it is still being supported, new designs are encouraged to use NX_DISABLE_IPV6_PURGE_UNUSED_CACHE_ENTRIES .
NX_IPV6_NEIGHBOR_CACHE_SIZE	Specifies the number of entries in the IPv6 Neighbor Cache table. Defined in <i>nx_nd_cache.h</i> , the default value is 16.
NX_MAX_MULTICAST_SOLICIT	Specifies the number of Neighbor Solicitation messages NetX Duo transmits as part of the IPv6 Neighbor Discovery protocol when mapping between IPv6 address and MAC address is required. Defined in <i>nx_nd_cache.h</i> , the default value is 3.
NX_MAX_UNICAST_SOLICIT	Specifies the number of Neighbor Solicitation messages NetX Duo transmits to determine a specific neighbor's reachability. Defined in <i>nx_nd_cache.h</i> , the default value is 3.
NX_ND_MAX_QUEUE_DEPTH	Symbol that defines the maximum number of packets queued up for ND cache to be resolved. By default this symbol is set to 4.

Option	Description
NX_REACHABLE_TIME	Specifies the time out in seconds for a cache entry to exist in the REACHABLE state with no packets received from the cache destination IPv6 address. Defined in <i>nx_nd_cache.h</i> , the default value is 30.
NX_RETRANS_TIMER	Specifies in milliseconds the length of delay between solicitation packets sent by NetX Duo. Defined in <i>nx_nd_cache.h</i> , the default value is 1000.
NXDUP_DISABLE_DAD	Renamed to <i>NX_DISABLE_IPV6_DAD</i> . Although it is still being supported, new designs are encouraged to use <i>NX_DISABLE_IPV6_DAD</i> .
NXDUP_DUP_ADDR_DETECT_TRANSITS	Renamed to <i>NX_IPV6_DAD_TRANSITS</i> . Although it is still being supported, new designs are encouraged to use <i>NX_IPV6_DAD_TRANSITS</i> .

Miscellaneous ICMPv6 Configuration Options

Option	Description
NX_DISABLE_ICMPV6_ERROR_MESSAGE	disables NetX Duo from sending an ICMPv6 error message in response to a problem packet (e.g., improperly formatted header or packet header type is deprecated) received from another host.
NX_DISABLE_ICMPV6_REDIRECT_PROCESS	disables ICMPv6 redirect packet processing. NetX Duo by default processes redirect messages and updates the destination table with next hop IP address information.
NX_DISABLE_ICMPV6_ROUTER_ADVERTISEMENT_PROCESS	processing information received in IPv6 router advertisement packets.

Option	Description
NX_DISABLE_ICMPV6_ROUTER_SOLICITATION	NetX Duo from sending IPv6 router solicitation messages at regular intervals to the router.
NX_DISABLE_ICMPV6_RX_CHECKSUM	Disabled, disables ICMPv6 checksum computation on received ICMP packets.
NX_DISABLE_ICMPV6_RX_CHECKSUM	Renamed to <i>NX_DISABLE_ICMPV6_RX_CHECKSUM</i> . Although it is still being supported, new designs are encouraged to use <i>NX_DISABLE_ICMPV6_RX_CHECKSUM</i> .
NX_DISABLE_ICMPV6_TX_CHECKSUM	Disabled, disables and ICMPv6 checksum computation on transmitted ICMP packets.
NX_DISABLE_ICMPV6_TX_CHECKSUM	Renamed to <i>NX_DISABLE_ICMPV6_TX_CHECKSUM</i> . Although it is still being supported, new designs are encouraged to use <i>NX_DISABLE_ICMPV6_TX_CHECKSUM</i> .
NX_ICMPV6_MAX_RTR_SOLICITATIONS	The max number of router solicitations a host sends until a router response is received. If no response is received, the host concludes no router is present. The default value is 3.
NX_ICMPV6_RTR_SOLICITATION_INTERVAL	The maximum delay for the initial router solicitation in seconds.
NX_ICMPV6_RTR_SOLICITATION_INTERVAL	The interval between two router solicitation messages. The default value is 4.
NXDUP_DESTINATION_TABLE_SIZE	Renamed to <i>NX_IPV6_DESTINATION_TABLE_SIZE</i> . Although it is still being supported, new designs are encouraged to use <i>NX_IPV6_DESTINATION_TABLE_SIZE</i> .
NXDUP_DISABLE_ICMPV6_ERROR_MESSAGE	Renamed to <i>NX_DISABLE_ICMPV6_ERROR_MESSAGE</i> . Although it is still being supported, new designs are encouraged to use <i>NX_DISABLE_ICMPV6_ERROR_MESSAGE</i> .

Option	Description
NXDUO_DISABLE_ICMPV6_REDIRECT_PROCESS	<i>NX_DISABLE_ICMPV6_REDIRECT_PROCESS.</i> Although it is still being supported, new designs are encouraged to use <i>NX_DISABLE_ICMPV6_REDIRECT_PROCESS</i>
NXDUO_DISABLE_ICMPV6_ROUTER_ADVERTISEMENT_PROCESS	<i>NX_DISABLE_ICMPV6_ROUTER_ADVERTISEMENT_PROCESS.</i> Although it is still being supported, new designs are encouraged to use <i>NX_DISABLE_ICMPV6_ROUTER_ADVERTISEMENT_PROCESS</i>
NXDUO_DISABLE_ICMPV6_ROUTER_SOLICITATION	<i>NX_DISABLE_ICMPV6_ROUTER_SOLICITATION.</i> Although it is still being supported, new designs are encouraged to use <i>NX_DISABLE_ICMPV6_ROUTER_SOLICITATION.</i>
NXDUO_ICMPV6_MAX_RTR_SOLICITATIONS	<i>NX_ICMPV6_MAX_RTR_SOLICITATIONS.</i> Although it is still being supported, new designs are encouraged to use <i>NX_ICMPV6_MAX_RTR_SOLICITATIONS.</i>
NXDUO_ICMPV6_RTR_SOLICITATION_INTERVAL	<i>NX_ICMPV6_RTR_SOLICITATION_INTERVAL.</i> This symbol is being depreciated. Although it is still being supported, new designs are encouraged to use <i>NX_ICMPV6_RTR_SOLICITATION_INTERVAL</i>

NetX Duo Version ID

The current version of NetX Duo is available to both the user and the application software during runtime. You can obtain the NetX Duo version from examination of the **nx_port.h** file. In addition, this file also contains a version history of the corresponding port. Application software can obtain the NetX Duo version by examining the global string **_nx_version_id** in **nx_port.h**.

Application software can also obtain release information from the constants shown below defined in **nx_api.h**.

These constants identify the current product release by name and the product major and minor version.

```
#define EL_PRODUCT_NETXDUA
#define NETXDUA_MAJOR_VERSION
#define NETXDUA_MINOR_VERSION
```